

# Deployment and Operations for Software Engineers 2<sup>nd</sup> Ed

## Chapter 2—Virtualization



# Outline

- **Virtual machines**
- Virtual machine images
- Containers
- Container repositories
- Serverless architecture




# What is a Virtual Machine?

- A Virtual Machine (VM) is a software construct that exposes the resources of CPU, memory, disk, and network connection to loaded software making it appear to the loaded software that it is executing on a physical computer.
- A VM runs under the control of a specialized operating system called a hypervisor.




# Type 1 Hypervisor



The hypervisor is the operating system for virtual machines  
A type 1 hypervisor runs on bare metal.  
C



# Type 2 Hypervisor



- A type 2 hypervisor runs on top of your laptop's operating system.
- From the perspective of the laptop's OS, it is an application.



# VM resources

- From the perspective of software executing inside a VM, it is the same as software executing on a bare metal computer.
- A VM has the same resources as a physical computer.
  - CPU
  - Memory
  - Disk
  - Network connection



# Discussion questions

1. How are the actions of one VM kept separate from the actions of another VM on the same host?  
(hint: think one resource at a time)
2. How does a VM write to a disk?
3. How does a VM write to the network?
4. How is a VM created?
5. How are VMs booted?



# Outline

- Virtual machines
- **Virtual machine images**
- Containers
- Container repositories
- Serverless architecture






# VM Image

- A VM image is a set of bits on a disk (or from over a network) that contains all of the software, directories, and data from its created VM
- An image is created by
  - building a VM or physical computer with all of the desired software.
  - Writing the contents of memory to a disk file (or sending it over a network)
- A VM is created by reading an image.



# Creating an image and using it to create another VM





# Image across network

- Suppose image is 8 Gb(yte).
- Reasonable network speed is 1 Gb(it)/sec (rated)
- 8 bits per byte
- Moving image, in theory, will take 64 seconds.
- In practice a 1Gb(it) network speed is .35 Gb(it)/sec
- Transferring image will take ~3 minutes.
- VM still must be booted to execute in new location.



# Discussion questions

- Can a VM image created on AWS be moved to Google Cloud?



# Outline

- Virtual machines
- Virtual machine images
- **Containers**
- Container repositories
- Serverless architecture

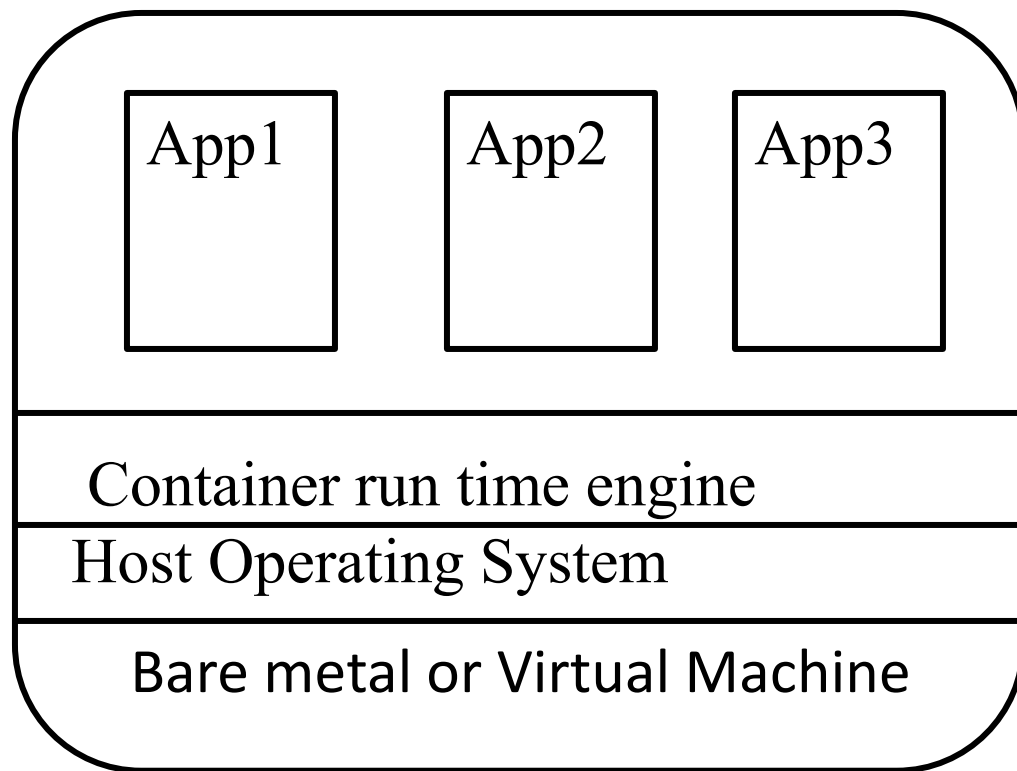


# Container is lighter weight virtualization mechanism

- Containers are a mechanism to maintain the advantages of virtualization— isolation of its actions from other containers—while reducing image transfer time and startup time.
- A container is an executable image that packages a service and its dependent libraries. It runs under the control of a container engine.
- Containers are isolated from each other in terms of
  - address space
  - Disk usage
  - Processor usage
  - Network usage (with some exceptions).



# Executing container



- Containers execute on top of a container run time engine
- This, in turn, executes on top of an operating system
- Means containers share a single operating system



# Container layers

- Containers can be segmented into layers.
  - E.g. consider the LAMP stack – Linux, Apache, MySQL, PHP
  - Each item can be made a separate layer
- Suppose only the top layer has changed. Then when moving the container, only the top layer must be moved. E.g., a new version of PHP in our example.
- This makes moving a container much lighter weight than moving a VM. On the order of milliseconds rather than minutes.





# Comparing VMs and containers

- VMs virtualize hardware
- Containers virtualize operating systems
- VMs take minutes to transfer over a network
- Containers can be transferred in milliseconds—depending on the mechanisms used
- VMs are monitored, started and stopped by the hypervisor
- Containers are monitored, started and stopped by the run time engine.



# Discussion questions

1. How are the actions of one container kept separate from the actions of another container on the same host? On the same runtime engine?
2. How does a container write to a disk?
3. How does a VM container write to the network?
4. How is a container created?
5. How are container placed into execution?
6. How does the container management system know that only the top layer has been changed?



# Outline

- Virtual machines
- Virtual machine images
- Containers
- **Container repositories**
- Serverless architecture



# Container repositories

- Container images are typically stored in repositories
- Similar to version control systems
  - Accessible with permissions
  - Push/pull interface
  - Images can be tagged with version numbers



# Docker Hub

- Docker Hub is a publicly available repository.
- Since it is open source, It is also used for private container repositories
- It is the most widely used container repository
- Private repositories can be access controlled



# Integrating with development workflow

- Multiple team members may wish to share images
- Images can be in production, under development or under test
- Private container repository allows images to be stored and shared.
  - Any image can be “pulled” to any host
  - Tagging as “latest” allows updates to be propagated. Pull <image name>:latest gets the last image checked into repository with that name.



# Discussion questions

1. What are container repositories other than Docker Hub? Why would you chose one over the other?



# Outline

- Virtual machines
- Virtual machine images
- Containers
- Container repositories
- **Serverless architecture**





# Serverless architecture

- A simple example using PHP.
- PHP is an interpreted language.
  - A file containing your program is fed to the interpreter to be executed.
- Cloud provider has pool of PHP servers.
  - You tell one of the pool the location of the file with your program.
  - Your program is executed immediately.



# General definition

- A serverless architecture consists of
  - A pool of servers with container run time engines
  - A stateless container with no dependencies on your other code
- Since the container has no dependencies, it can be executed directly by run time engine.
- Container is deployed and activated when a message arrives for it.
- Loads in milliseconds depending on cloud provider's caching strategy.
- For AWS, only one request per instance.



# More complicated example

- Suppose you wish to send an email welcome message to a new subscriber.
- Your container
  - Consists of code to prepare and end the email.
  - Parameters are email address and subscriber's name.
- Container is stateless. Necessary information comes from request parameters.
- Your full system sends request to the welcome message container
  - This triggers the cloud provider to load and activate it.
- Email message is sent.



# Discussion questions

1. How does the cloud provider know that a container is intended to be serverless?
2. What are restrictions on serverless container other than it must be stateless?